

Tree-Structured Indexes (Brass Tacks)


Chapter 10 Ramakrishnan and Gehrke
(Sections 10.3-10.8)

What will I learn from this set of lectures?

- ❖ How do B+trees work (for search)?
- ❖ How can I tune B+trees for performance?
- ❖ How can I maintain its balance against inserts and deletes?
- ❖ How do I build a B+tree from scratch?
- ❖ How can I handle key values that are very long - e.g., long song names or long names of people?


B+ Tree: The Most Widely Used Index

- ❖ Insert/delete at $\log_F N$ cost; keep tree *height-balanced*. (F = fanout, N = # leaf pages)
- ❖ Minimum 50% occupancy (except for root). Each node contains $d \leq \underline{m} \leq 2d$ entries. The parameter d is called the *order* of the tree.
- ❖ Supports equality and range-searches efficiently.




Example B+ Tree

- ❖ Search begins at root, and key comparisons direct it to a leaf (as in ISAM)
- ❖ Binary search within a node
- ❖ Search for 5*, 15*, all data entries $\geq 28^*$



insert

B+ Trees basics



The key 50,000 separates or discriminates between N and its sibling. Plays a crucial role in search and update maintenance. Call 50,000 the **separator/discriminator** of N and its sibling.

B+ Trees in Practice

- ❖ Typical order: 100. Typical fill-factor: $\ln 2 = 66.5\%$ (approx).
 - average fanout = $2 \times 100 \times 66.5\% = 133$
- ❖ Typical capacities:
 - Height 4: $133^4 = 312,900,721$ pages.
 - Height 3: $133^3 = 2,352,637$ pages
- ❖ Can often hold top levels in buffer pool:
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 MByte (approx.)
 - Level 3 = 17,689 pages = 133 MBytes (approx.)
 - Level 4 = 2,352,637 pages = 17.689 GBytes (approx.)
 - Level 5 = 312,900,721 pages = 2.352637 tera bytes! (approx.)
- ❖ For typical orders ($d \sim 100-200$), a shallow B+tree can accommodate very large files. **How tall a B+tree do we need to cover all of Canada's taxpayer records?**

B+ Trees in practice



- ❖ Suppose a node is implemented as a block, a block is 4 K, a key is 12 bytes, whereas a pointer is 8 bytes.
 - For a file occupying b (logical) blocks, what is the min/max/avg height of a B+tree index?
 - If we have m bytes of RAM, how many levels of the B+tree can we prefetch to speed up performance?

Inserting a Data Entry into a B+ Tree


- ❖ Find correct leaf L .
- ❖ Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must *split* L (into L and a new node $L2$)
 - ◆ Redistribute entries evenly, *copy up* middle key.
 - ◆ Insert index entry pointing to $L2$ (i.e., the middle key copied up) into parent of L .
- ❖ This can happen recursively
 - To *split index node*, redistribute entries evenly, but *bump up* middle key. (Contrast with leaf splits.)
- ❖ Splits “grow” tree; root split increases height.
 - Tree growth: gets *wider*, maybe even *one level taller*.

Inserting 8^* into Example B+ Tree

- ❖ Observe how minimum occupancy is guaranteed in both leaf and index pg splits.
- ❖ Note difference between *copy-up* and *bump-up*: Why do we handle leaf page split and index page split differently?



Example B+ Tree After Inserting 8*



- ❖ Notice that root was split, leading to increase in height.
- ❖ In this example, we can avoid split by re-distributing entries; however, this is usually not done in practice.

What would it look like?

delete


Deleting a Data Entry from a B+ Tree

- ❖ Start at root, find leaf L where entry belongs.
- ❖ Remove the entry.
 - If L is at least half-full, *done!*
 - If L has only $d-1$ entries,
 - ◆ Try to **re-distribute**, borrowing from *sibling* (adjacent node with same parent as L). Adjust the key that separates L and its sibling.
 - ◆ If re-distribution fails, **merge** L and sibling.
- ❖ If merge occurred, must delete separator entry (discriminating L & its sibling) from parent of L .
- ❖ Merge could propagate to root, decreasing height.

Question: Can L 's occupancy ever drop below $d-1$?

Example Tree After (Inserting 8* Then) Deleting 19* and 20*



...



- ❖ Deleting 19* is easy.
- ❖ Deleting 20* is done with re-distribution.
Notice how **new** middle key is **copied up**.


... And Then Deleting 24*

- ❖ Must merge.
- ❖ Observe '*toss*' of index entry (on right), and '*pull down*' of index entry (below).




Example of Non-leaf Redistribution

- ❖ Tree is shown below *during deletion* of 24*. (What could be a possible initial tree?)
- ❖ In contrast to previous example, can redistribute entry from left child of root to right child.



After Re-distribution

- ❖ Intuitively, entries are **re-distributed by "pushing through"** the splitting entry (i.e., separator) in the parent node.
- ❖ It suffices to re-distribute index entry with key 20; we've re-distributed 17 as well for illustration.



Optimization 1: Prefix Key Compression


- ❖ Important to increase fan-out. (Why?)
- ❖ Key values in index entries only 'direct traffic'; can often compress them.
 - E.g., If we have adjacent **index entries** with search key values *Dannon Yogurt*, *David Smith* and *Devarakonda Murthy*, we can abbreviate David *Smith* to *Dav*. (The other keys can be compressed too ...)
 - ◆ Is this correct? Not quite! What if there is a **data entry** *Davey Jones*? (Can only compress David *Smith* to *Davi*)
 - ◆ In general, while compressing, must leave each index entry greater than every key value (in any subtree) to its left.
- ❖ Insert/delete must be suitably modified.
- ❖ This idea works for any field/attribute whose data type is a long string: e.g., customer code, shipping tracking number, etc.

Analyzing insert/delete time.


- ❖ Given a B+tree index with height h , what can you say about best case, average case, and worst case I/O?
- ❖ It must have something to do with h .
- ❖ But is it exactly h ?
- ❖ What exactly constitutes best, average, and worst case?

Optimization 2: Bulk Loading of a B+ Tree


- ❖ If we have a large collection of records, and we want to create a B+ tree on some field, doing so by repeatedly inserting records is very slow.
- ❖ *Bulk Loading* can be done much more efficiently.
- ❖ *Initialization*: Sort all data entries, insert pointer to first (leaf) page in a new (root) page.




Bulk Loading (contd.)




Bulk Loading (contd.)




Bulk Loading (contd.)



Bulk Loading (contd.)



Bulk Loading (contd.)

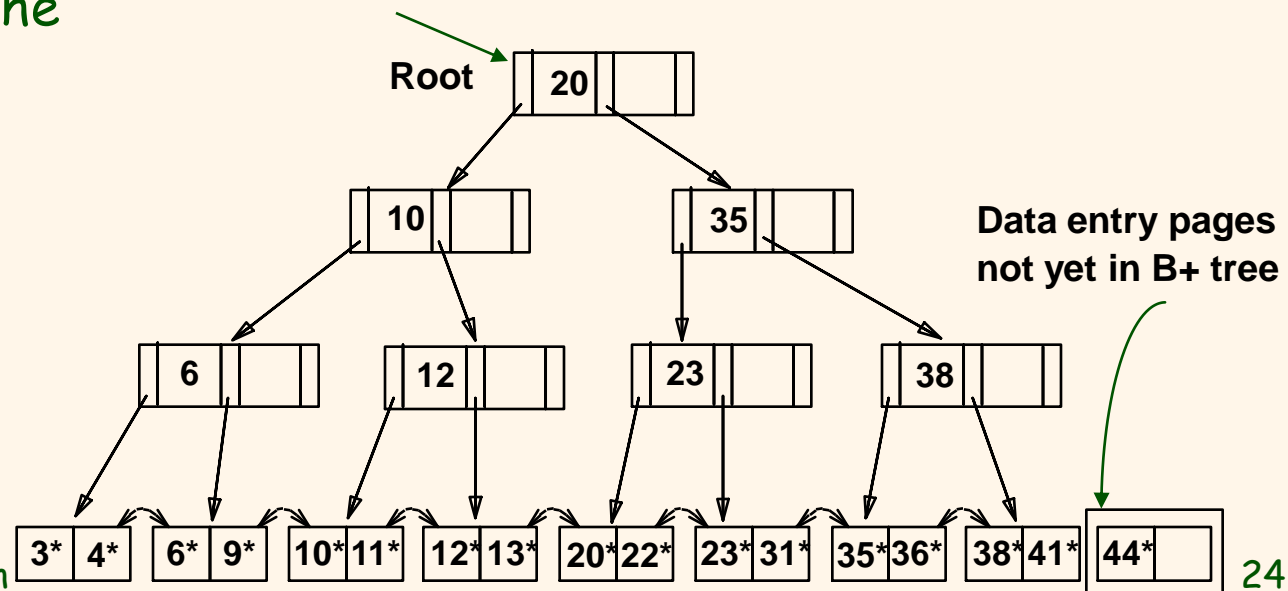



Bulk Loading (Contd.)

- ❖ Index entries for leaf pages always entered into right-most index page just above leaf level. When this fills up, it splits.

(Split may go up right-most path to the root.)

- ❖ Much faster than repeated inserts, especially when one considers locking!



Bulk Loading

- ❖ How would you estimate time taken to build a B+tree index from scratch?
 - Naïve approach of repeated insertions.
 - Bulk loading.
- ❖ Naïve approach: each insert is potentially a random block seek; cannot read/insert next data block until after previous block has been inserted. Very slow.
- ❖ Bulk loading: dominant factor - sorting data (or data entries as appropriate). Followed by another sequential read (w/ proper buffer management).
 - Explored in exercises.

Summary of Bulk Loading

- ❖ Option 1: multiple inserts.
 - Slow.
 - Does not ensure sequential storage of leaves.
- ❖ Option 2: bulk loading
 - Has advantages for concurrency control.
 - Fewer I/Os during build.
 - Leaves will be stored sequentially (and linked, of course).
 - Can control “fill factor” on pages.

Summary

- ❖ Tree-structured indexes are ideal for range-searches, also good for equality searches.
- ❖ ISAM is a static structure.
 - Only leaf pages modified; overflow pages needed.
 - Overflow chains can degrade performance unless size of data set and data distribution stay constant.
- ❖ B+ tree is a dynamic structure.
 - Inserts/deletes leave tree height-balanced; $\log_F N$ cost.
 - High fanout (F) means depth (i.e., height) rarely more than 3 or 4.
 - Almost always better than simply maintaining a sorted file.

Summary (Contd.)

- ❖ Key compression increases fanout, reduces height.
- ❖ Bulk loading can be much faster than repeated inserts for creating a B+ tree on a large data set.
- ❖ Most widely used index in database management systems because of its versatility. One of the most optimized components of a DBMS.